



B+ TREES

A B+ Tree is a B tree variation. Its terminal nodes or leaves are used to store records and nonterminals. Nonterminal nodes are used to store key values.





- Key size \leq Record size, so the B+ Trees are shallower (shorter) than B trees.
- In B+ trees all records require same disk accesses.
- All records are stored at the same level.
- Terminal nodes are connected each other.



B+ Tree Characteristics

- Dynamic: self organizing
- Balanced at all times.
- Supports an ordering on a large file
- Provides direct access based on a key value.
- Provides efficient range query processing
- Display st major when
 $St_name > Ali$ and $st_name > HUSEYIN$





B+ TREE EXAMPLE


Records to be inserted to the B+ tree;

80,50,100,90,60,65,70,75,55,64,51,76,77,78,200,300,150


Index node degree =2 (4 key values)

Data nodes degree =1 (2 records)


INSERT 80




INSERT 50




INSERT 100, 90




INSERT 60=> first node becomes 50, 60, 80. So the middle one 60 goes up




INSERT 65




INSERT 70=>second node becomes 65,70,80. So the middle one 70 goes up




INSERT 75




INSERT 55=> First node becomes 50,55,60. so The middle one 55 goes up.




INSERT 64=>third node becomes 64, 65, 70, so the middle one 65 goes up. Index node becomes 55,60,65,70,80. so the middle one 65 has to go up one level




INSERT 51, 76




INSERT 77




INSERT 78




INSERT 200=> Last node becomes 90,100, 200. so 100 goes up. Index node becomes 70,76,78,80 and 100. so the middle one 78 has to go up



INSERT 300




INSERT 150



A B+ tree of degree d (order d) has the following features:

1. In index nodes, no of keys $\geq d$ and $\leq 2d$
(root is a special case, see the previous example)
2. All internal nodes with k keys has $(k+1)$ children.
(root may not obey to this rule according to our example)



In the B+ tree two disk accesses are enough


It is shown that each node is $\ln 2$ (0.70, 70% full). Shown by Yoo, 1978.

pointer size = 4 bytes
key size = 8 bytes
12 bytes

index node size = 2400 bytes

$2400 / 12 = 200$ $d = 100$

A full node can contain 200 key values.
A typical contains $0.7 * 200 = 140$ keys



level	no of nodes	total size for all nodes
1	1	$1 * 2400 = 2400$ bytes
2	140	$140 * 2400 = 326000$
3	19600	$19600 * 2400 \approx 47000000$ bytes = 47 MB

key top two levels in main memory

first disk access

2nd disk access → data node


no of records in data nodes

$$\text{node size} / \text{record size} = 2400 / 240 = 10$$

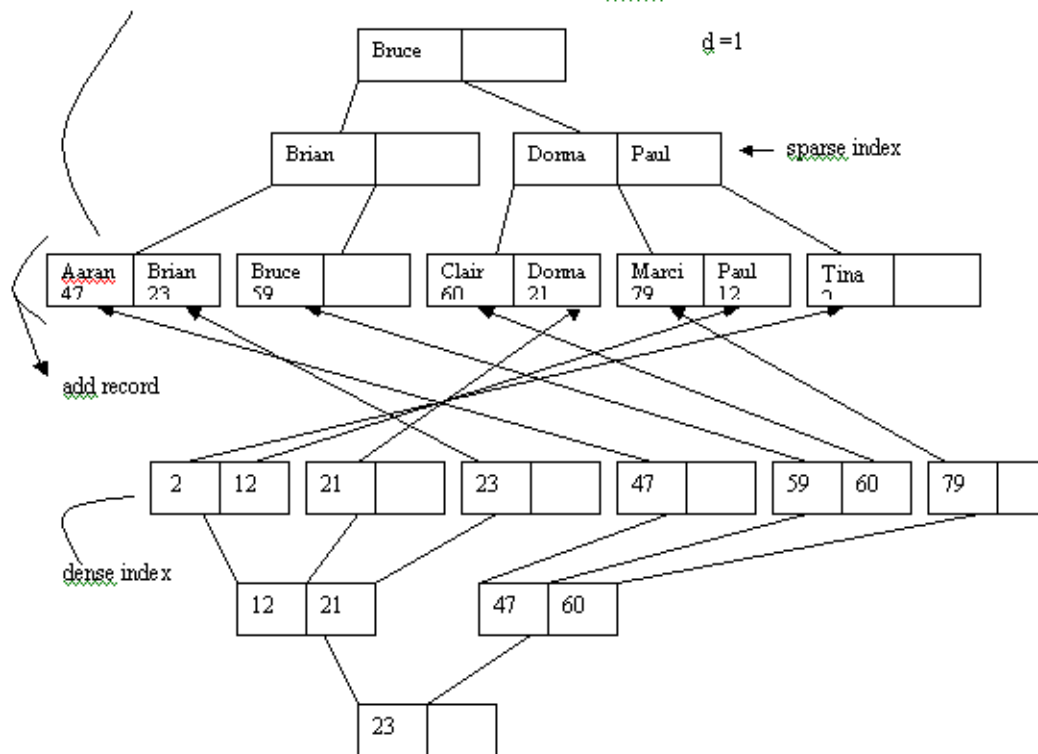
$$(2744000) * 0.70 * 10 \approx 19000000 \text{ records}$$

no of data nodes expected failures

Primary and Secondary Key Based Retrieval Using B+ trees





The order of records matches the order of index elements. (index is clustered)



- With 2 disk accesses how large a file can be handled by a B+ tree.

Assumption: 10 MB of memory is available for the index tree structure.



$bk / 140 = \text{number of index blocks at the parent of leaf nodes}$

$10 \text{ MB} \rightarrow 10,000,000 / 2400 = 4167 \text{ buckets}$
 bucket size in bytes

$$bk / (140)^2 + bk / (140)^2 + 1 = 4167$$

$$bk = 140^2 \times 4167$$

↑
number of available buckets in memory

|

fo: fan out : (expected number of pointers) / (index node)

$$bk = fo^2 \times mem \rightarrow mem = bk / (fo)^2$$

EXERCISE

Find the expected file size

From the calculations $bk = bk = 140^2 \times 4167$ we obtain 80 million buckets. If each data bucket is one block, 80 million buckets, file size is $2 \times 10^8 = 200000MB$.

One disk drive for IBM 3380 is 2400 MB so we get 80 disks as the expected size of the file in IBM 3380.

SIMPLE FILE OPERATIONS USING B+ TREES


n : number of nodes

$Bkfr$: number of records/bucket

Bk : No of data buckets = $n / \ln 2 * Bkfr$ ($\ln 2 = 0.7$)


- Time to read the whole file
 $T_x(\text{primary}) = bk * (s + r + dtt)$

- Finding the next record
 $TN(\text{primary}) = (1 / \ln 2 * Bkfr) * (s + r + dtt)$
This is the expected time
Min time needed = 0;
Max time needed = $(s + r + dtt)$




FILE INTERSECTION

F1



∩

F2



= RESULTANT FILE
(70000 records)

$n=100,000$

$n=100,000$

70% common records.

Record size=400 bytes

File size = 40MB

To create the intersection file look pile, compare in F2 for each two disk accesses.

Time needed to read F1 (Assume that we have 10MB memory) = $b \cdot ebt$

Block size = 2400 bytes

$Bfr = 6$

$b = 100,000 / 6 = 16667 \Rightarrow b \cdot ebt = 16667 \cdot 0.84 = 14 \text{ sec}$

Writing the resultant file

$70,000 / 6 = 11667 \text{ blocks}$

$11667 \cdot 0.84 = 10 \text{ sec}$

Searching the file

2 disk accesses / record

$= 2 \cdot n \cdot (s + t + btt)$


$= 2 \cdot 100,000 \cdot (s + t + btt) = 84 \text{ min}$

- Using two sorted files is much shorter = 28 sec

SECONDARY KEY BASED RETRIEVAL

Display employ-name when


Dept TOY and Age = 25



Conjunctive query:and....and....
Disjunctive query:or.....or.....
Mixed query: (and) or....

INVERTED FILE STRUCTURE

<u>Emp no</u>	<u>Emp name</u>	<u>dept</u>	<u>No of dependants</u>	<u>Salary</u>
01	Jane	Gun	1	
02	George	Shoe	2	
04	Dane	Shoe	3	
05	Mary	Gun	2	
06	Mike	Shoe	1	



Inverted indexes for secondary attributes

<u>Dept</u>	<u>Posting list</u>
GUN	01 05
SHOE	02 04 06

EXAMPLES ABOUT THIS WEEK LECTURE NOTES

- There are 8 million records. 400 bytes each records. These files are organized in a B+ tree.

- How much space does this file takes if $\ln 2$ of the space is used in the tree.

8 million * 400 bytes = 3200 MB is the file size. As only $\ln 2$ of the space is used 4500 MB of space is needed.

- Assume that the leaves are 2400 bytes. How many bytes is the parent of leaf level. (totally a bucket is 12 byte for an address and pointer)


The number of buckets is $192\,000\,000 \left((8\text{ million} * 400 / 2400) / \ln 2 \right)$

Each of these buckets needs 12 bytes for an address and pointer. ,


$1\,920\,000 * 12 = 22.8\text{ MB}$ is needed for the total buckets.

But only $\ln 2$ space is used so we have to divide this number with $\ln 2$ in order to get the actual space needed at parent of leaf level. $22.8 / \ln 2 = 33\text{ MB}$ is the parent of leaf level

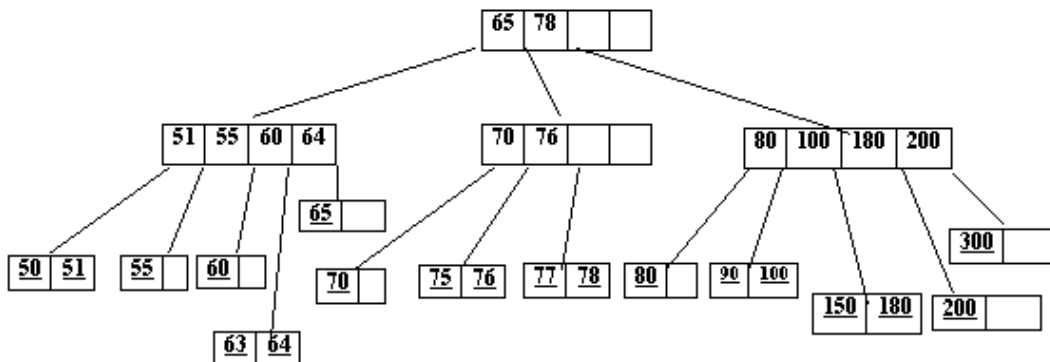
2. A Final B+ tree structure is given below. Insert new records with given key values **63**, **180**, **95**.



Insert 63=> In data field 4th node becomes 63,64,65. so the middle one 64 goes up and tree is structured again.



Insert 180=> In the data field 11th node becomes 150, 180, 200. so the middle one 180 goes up. Tree is restructured



Insert 95=> In the data field 10th node becomes 90, 95, 100. So the middle one 95 goes up. But in the index field 3rd node is full so the middle index 100 goes up one level.

